



Deep learning Models for Image/Video

Prof. Venki Muthukumar

What is OpenCV

- **OpenCV: Open-Source Computer Vision & Machine Learning software library**
 - Created in 1999 by Intel
 - Supported from 2008 by Willow Garage(a SME dedicated to hardware & open-source software for personal robotics applications)
www.willowgarage.com/pages/software/overview
 - Willow Garage also supports the Point Cloud Library (PCL)
 - **Cross-platform**
 - Windows | Linux | Android | MacOS | iOS*
 - **Free under BSD License**
 - Commercial & non-commercial applications



OpenCV

- Written in C / C++
- Extended to Python
- Stable source code opencv.org/releases
- Developments
 - Source code github.com/opencv
 - Wiki github.com/opencv/opencv/wiki
- APIs available for a variety of programming languages




- Current stable version is **4.13.0**

OpenCV Overview:

> 2000 algorithms

Robot support



opencv.willowgarage.com




General Image Processing Functions




Segmentation




Geometric descriptors



Features





Transforms



Machine Learning:

- Detection,
- Recognition



Matrix Math









Image Pyramids




Camera calibration, Stereo, 3D

Utilities and Data Structures




Fitting





Where is OpenCV Used?

- Google Maps, Google street view, Google Earth, Books
- Academic and Industry Research
- Safety monitoring (Dam sites, mines, swimming pools)
- Security systems
- Image retrieval
- Video search
- Structure from motion in movies
- Machine vision factory production inspection systems
- Robotics



OpenCV & DL

- Tensorflow, Pytorch, Caffe, MxNet, CNTK, Keras etc. are the libraries use OpenCV to build deep learning vision models.






OpenCV Resources

- **Online documentation** docs.opencv.org
- Reference | Tutorials | QuickStart | Examples
- **Online resources** [opencv.org>Resources](https://opencv.org/Resources)
- Books | Publications | Usefullinks
- **Q&A Forum** answers.opencv.org




Bibliography




OpenCV Essentials

Acquire, process, and analyze visual content to build full-fledged imaging applications using OpenCV




OpenCV Computer Vision Application Programming Cookbook Second Edition

Over 50 recipes to help you build computer vision applications in C++ using the OpenCV library




INSTANT OpenCV Starter

Get started with OpenCV using practical, hands-on projects




Learning Image Processing with OpenCV

Exploit the amazing features of OpenCV to create powerful image processing applications through easy-to-follow examples



Practical OpenCV

HANDS ON PROJECTS FOR COMPUTER VISION ON THE WINDOWS, LINUX AND RASPBERRY PI PLATFORMS




A Practical Introduction to COMPUTER VISION WITH OPENCV

WILEY



Mastering OpenCV with Practical Computer Vision Projects

Step-by-step tutorials to solve common real-world computer vision problems for desktop or mobile, from augmented reality and marker-based recognition to face recognition and 3D head tracking



Learning OpenCV 3

COMPUTER VISION IN C++ WITH THE OPENCV LIBRARY









OpenCV Blueprints

Expand your knowledge of computer vision by building amazing projects with OpenCV 3

Bibliography





Python | Java | .NET APIs



Bibliography

Mobile & Game APIs

Android | iOS | Raspberry Pi | Unity





OpenCV Modules

- **Calib3d**
 - Calibration, stereo, homography, rectify, projection, solvePnP
- **Contrib**
 - Octree, self-similar feature, sparse L-M, bundle adj, chamfer match
- **Core**
 - Data structures, access, matrix ops, basic image operations
- **features2D**
 - Feature detectors, descriptors and matchers in one architecture
- **Flann** (Fast library for approximate nearest neighbors)
- **Gpu** – CUDA speedups
- **Highgui**
 - Gui to read, write, draw, print and interact with images
- **Imgproc** – image processing functions
- **ML** – statistical machine learning, boosting, clustering
- **Objdetect** – PASCAL VOC latent SVM and data reading
- **Traincascade** – boosted rejection cascade



Digital Image Acquisition




FIGURE 2.15 An example of the digital image acquisition process. (a) Energy (“illumination”) source. (b) An element of a scene. (c) Imaging system. (d) Projection of the scene onto the image plane. (e) Digitized image.


Digital Camera Issues

- Noise
 - caused by low light
- Color
 - color fringing (chromatic aberration) artifacts from Bayer patterns
- Blooming
 - charge overflowing into neighboring pixels
- In-camera processing
 - over-sharpening can produce halos
- Compression
 - creates blocking artefacts





Image Filtering



Input Image

7	8	4	5	5
5	9	4	3	8
5	2	7	2	2
6	1	9	2	4
3	2	6	9	4

median

2	2	2	2	4	4	4	9	9
---	---	---	---	---	---	---	---	---


Filtered Img

8	5	4	4	5
7	5	4	4	3
5	5	3	4	2
2	5	2	4	2
2	6	2	4	4

FIGURE 3.28 The mechanics of linear spatial filtering using a 3×3 filter mask. The form chosen to denote the coordinates of the filter mask coefficients simplifies writing expressions for linear filtering.



Application of Filtering




a b c

FIGURE 3.34 (a) Image of size 528×485 pixels from the Hubble Space Telescope. (b) Image filtered with a 15×15 averaging mask. (c) Result of thresholding (b). (Original image courtesy of NASA.)

Filtering for Denoising





Embedded ML



a b c

FIGURE 3.35 (a) X-ray image of circuit board corrupted by salt-and-pepper noise. (b) Noise reduction with a 3×3 averaging mask. (c) Noise reduction with a 3×3 median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)

Edge detection



- Convert a 2D image into a set of curves
 - Extracts salient features of the scene
 - More compact than pixels

Optimal Edge Detection: Canny



Embedded ML

- Assume:
 - Linear filtering
 - Additive Gaussian noise
- Edge detector should have:
 - Good Detection. Filter responds to edge, not noise.
 - Good Localization: detected edge near true edge.
 - Single Response: one per edge.
- Detection/Localization trade-off
 - More smoothing improves detection
 - And hurts localization.




The Canny edge detector



original image (Lena)



Norm of gradient



thresholding



thining




Contours Examples



Source Picture
 $(300 \times 600 = 180000 \text{ pts total})$



Retrieved Contours
($<1800 \text{ pts total}$)




After Approximation
($<180 \text{ pts total}$)

And it is rather fast: $\sim 70 \text{ FPS}$ for 640×480 on complex scenes

Object Recognition




Embedded ML



Human Pose/Sign Estimation



Embedded ML






Motion Detection

- **A. Frame Differencing (Basic)**
 - *Logic:* Subtracts the current frame from the previous one ($|Frame_t - Frame_{t-1}|$).
 - *Pros/Cons:* Extremely fast, but sensitive to noise and fails if the object stops moving.
 - *Key Function:* `cv2.absdiff(frame1, frame2)`
- **B. Background Subtraction (Intermediate)**
 - *Logic:* The system "learns" a static background model over time and isolates moving foreground objects. Adapts to lighting changes.
 - *Pros/Cons:* Robust for surveillance; handles shadows.
 - *Key Function:* `cv2.createBackgroundSubtractorMOG2()`
- **C. Optical Flow (Advanced)**
 - *Logic:* Tracks the apparent motion of brightness patterns to calculate velocity and direction vectors for every pixel.
 - *Pros/Cons:* Provides speed and direction data, but is computationally expensive.
 - *Key Function:* `cv2.calcOpticalFlowFarneback()`

Optical Flow

```
// opencv/samples/c/lkdemo.c
int main(...){
    ...
    CvCapture* capture = <...> ?
        cvCaptureFromCAM(camera_id) :
        cvCaptureFromFile(path);
    if( !capture ) return -1;
    for(;;) {
        IplImage* frame=cvQueryFrame(capture);
        if(!frame) break;
        // ... copy and process image
        cvCalcOpticalFlowPyrLK( ... )
        cvShowImage( "LkDemo", result );
    }
}
```


calcOpticalFlowPyrLK()
*Also see dense optical flow:
calcOpticalFlowFarneback()*



$$I(x + dx, y + dy, t + dt) = I(x, y, t);$$
$$-\partial I / \partial t = \partial I / \partial x \cdot (dx / dt) + \partial I / \partial y \cdot (dy / dt);$$

$$G \cdot \partial X = b,$$

$$\partial X = (\partial x, \partial y), G = \begin{bmatrix} I_x^2, & I_x I_y \\ I_x I_y, & I_y^2 \end{bmatrix}, b = \sum I_t \begin{bmatrix} I_x \\ I_y \end{bmatrix}$$




Single Camera Calibration

See samples/cpp/calibration.cpp

Now, camera calibration can be done by holding checkerboard in front of the camera for a few seconds.

And after that you'll get:

3D view of checkerboard



Un-distorted image



Machine Learning Library (MLL)

CLASSIFICATION / REGRESSION

(new) **Fast Approximate NN (FLANN)**

(new) **Extremely Random Trees**

(coming) **LSH**

CART

Naïve Bayes

MLP (Back propagation)

Statistical Boosting, 4 flavors

Random Forests

SVM

Face Detector

(Histogram matching)


(Correlation)

CLUSTERING

K-Means

EM

(Mahalanobis distance)




TUNING/VALIDATION

Cross validation

Bootstrapping

Variable importance

Sampling methods



YOLO

- **Definition:** YOLO is a state-of-the-art, real-time object detection system.
- **The Core Concept:** Unlike traditional classifiers that scan images step-by-step, YOLO treats object detection as a single regression problem.
- **"You Only Look Once":** The model looks at the entire image exactly once to predict:
 - What objects are present (Classification).
 - Where they are located (Localization/Bounding Boxes).
- **Key Characteristic:** Extremely fast inference speed, making it the standard for video processing and live streams.



Evolution of YOLO

- **YOLOv1 (2015)**: The breakthrough. Fast but struggled with small objects.
- **YOLOv2 & v3**: Introduced "Anchor Boxes" and multi-scale detection (detecting small & big objects simultaneously).
- **YOLOv4 & v5**: Focused on speed optimizations (Mosaic augmentation) and usability (PyTorch implementation).
- **YOLOv8 (2023) & Beyond**: The current standard. Introduced "Anchor-Free" detection, simpler architecture, and native support for segmentation and pose estimation.
- **YOLO-NAS / YOLO-World**: New variants using Neural Architecture Search and open-vocabulary detection.
-

Pros & Cons of YOLO

- **Advantages:**
- **Speed:** Capable of 140+ FPS (frames per second) on modern GPUs.
- **Global Context:** Because it sees the whole image at once, it makes fewer "background errors" compared to sliding-window models.
- **Versatility:** Can perform Detection, Segmentation (masks), and Pose Estimation (skeletons) with the same API.
- **Limitations:**
- Can struggle with very small objects grouped tightly together (e.g., a flock of birds) compared to slower, two-stage detectors like Faster R-CNN.

YOLO Real-World Applications



- **Autonomous Driving:** Detecting pedestrians, traffic lights, and other cars in milliseconds.
- **Security & Surveillance:** Intruder detection, face recognition, and crowd counting.
- **Retail:** Automated checkout (identifying products) and people counting in stores.
- **Healthcare:** Detecting tumors in X-rays or counting cells in microscopy.
- **Agriculture:** Weed detection and crop monitoring from drones.



Embedded ML